

# **Transformer Architecture Implementation Details**

Forward & Backward Pass Analysis

with Tensor Parallelism

November 3, 2025

# Contents

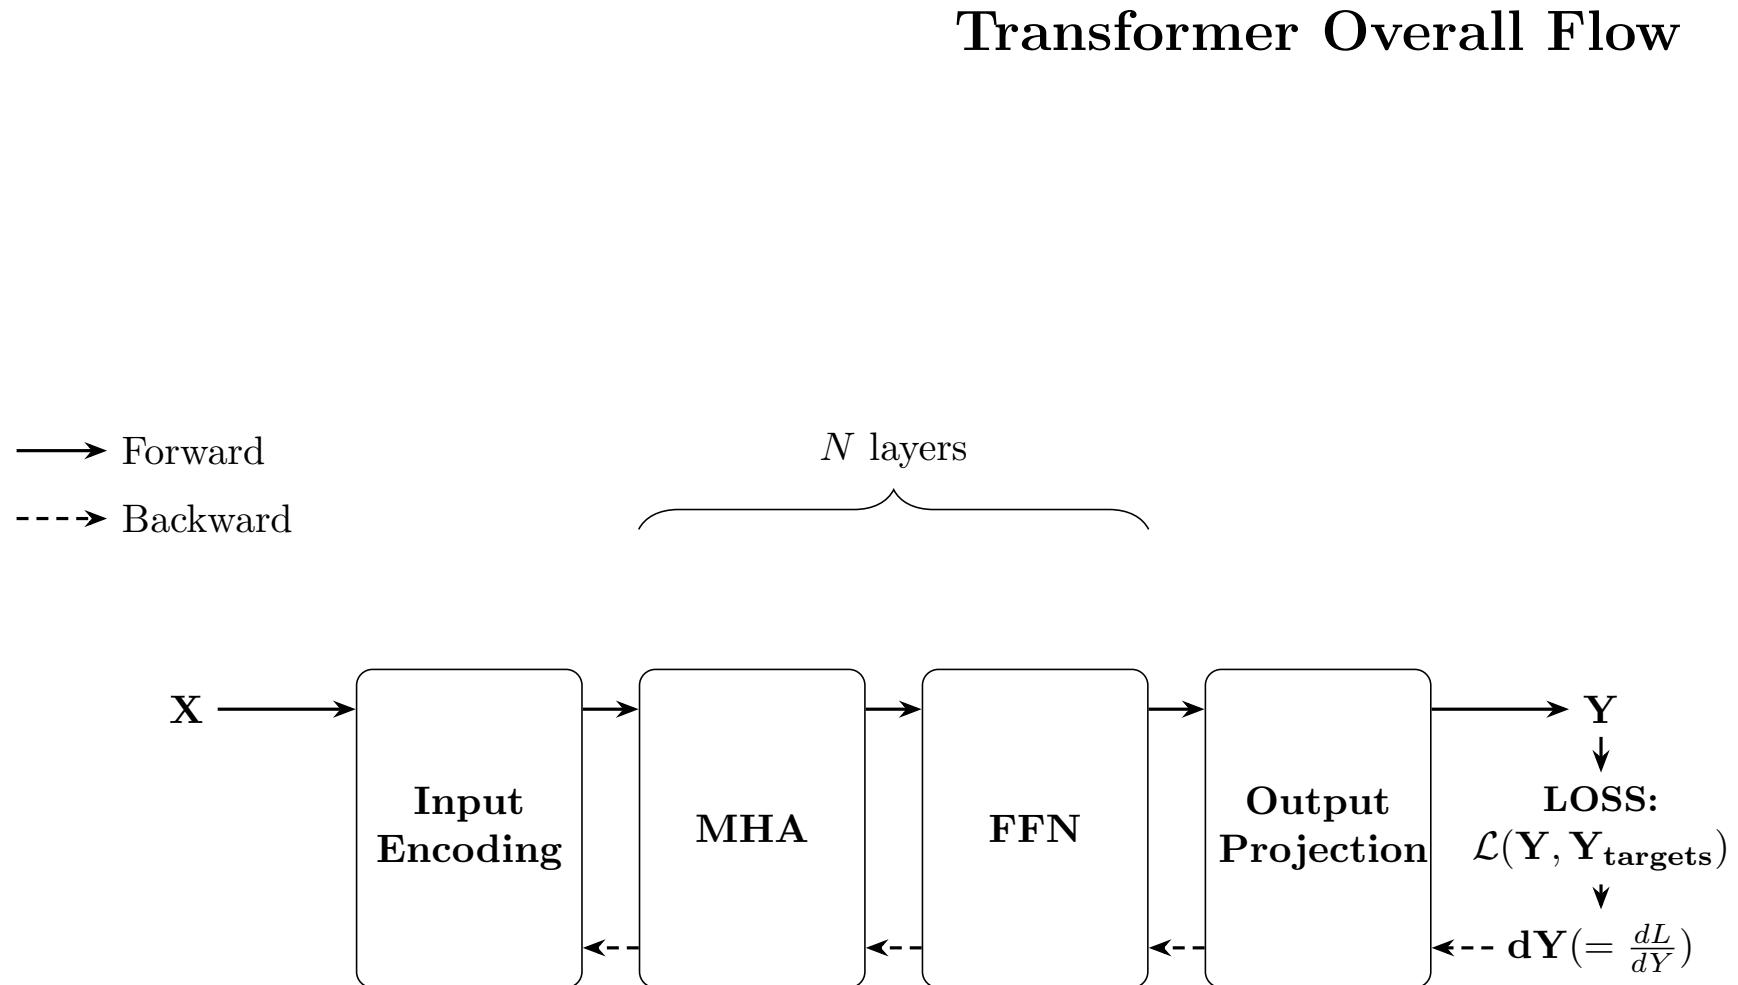
<b>1 Single Node Transformer</b>	<b>3</b>
1.1 Overall Architecture & Data Flow . . . . .	4
1.2 Input Embedding Layer . . . . .	5
1.2.1 Forward Pass . . . . .	5
1.2.2 Operations Summary . . . . .	6
1.3 Multi-Head Attention (MHA) . . . . .	7
1.3.1 Forward Pass . . . . .	7
1.3.2 Backward Pass . . . . .	8
1.3.3 Operations Summary . . . . .	10
1.4 Feed-Forward Network (FFN/MLP) . . . . .	11
1.4.1 Forward Pass . . . . .	11
1.4.2 Backward Pass . . . . .	12
1.4.3 Operations Summary . . . . .	13
1.5 Output Projection & Loss Computation . . . . .	14
1.5.1 Forward Pass . . . . .	14
1.5.2 Backward Pass . . . . .	15
1.5.3 Operations Summary . . . . .	16
<b>2 Tensor Parallelism (TP)</b>	<b>17</b>
2.1 TP Overview . . . . .	18
2.2 Forward Pass . . . . .	19
2.3 Backward Pass . . . . .	20
2.4 Communication Patterns . . . . .	20
<b>3 Data Parallelism (DP)</b>	<b>22</b>
3.1 DP Overview . . . . .	22

## 1 Single Node Transformer

This chapter covers the complete forward and backward pass of a Transformer model running on a single node (no parallelism).

## 1.1 Overall Architecture & Data Flow

The following diagram shows the high-level architecture of a Transformer layer, including both forward and backward passes.

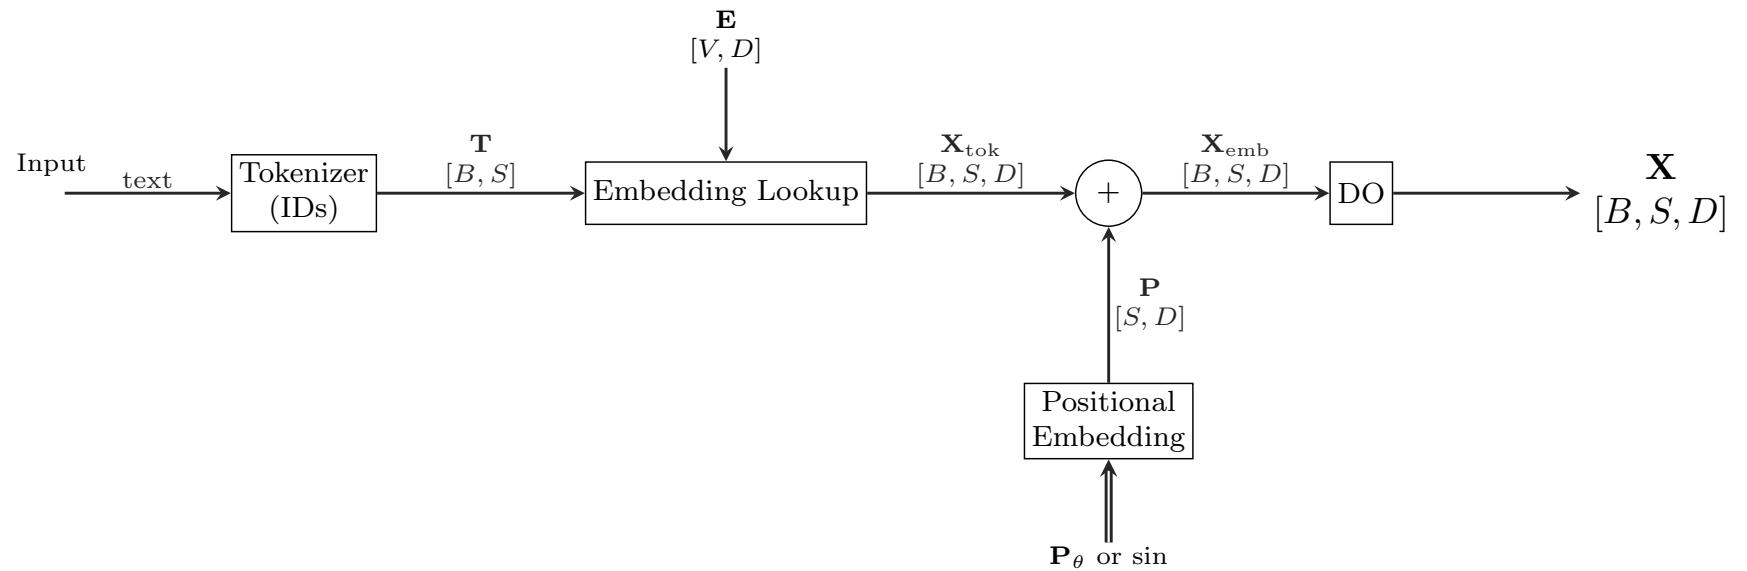


## 1.2 Input Embedding Layer

The input embedding layer converts token indices into dense vector representations and adds positional encodings.

### 1.2.1 Forward Pass

**Input → Embedding → LN (Input to MHA)**



### 1.2.2 Operations Summary

Operations (Ops)			
Abbrev	Name	Type / Shape	Notes
Tokenizer	Tokenizer (IDs)	op	Maps raw text → integer ids $\mathbf{T} \in \mathbb{Z}^{[B,S]}$ .
Embedding Lookup	Embedding Lookup	op	Gathers rows from $\mathbf{E} \in \mathbb{R}^{V \times D}$ using ids $\mathbf{T}$ .
+	Element-wise Add (dashed circle)	op	Adds token and positional embeddings; broadcasting over $B, S$ if needed.
DO	Dropout	op	Training-time stochastic dropout on $\mathbf{X}_{\text{emb}}$ ; identity at inference.
(none)	Broadcast $\text{BC}_{B,S}(\cdot)$	op	Expands $[S, D]$ (or $[D]$ ) to $[B, S, D]$ across batch/sequence.

Data Tensors (Values)			
Symbol	Name	Shape	Notes
text	Raw input text	—	Character/byte stream before tokenization.
$\mathbf{T}$	Token ids	$[B, S]$	Output of Tokenizer; integers in $\{0, \dots, V-1\}$ .
$\mathbf{E}$	Embedding matrix (params)	$[V, D]$	Trainable; each vocab entry has a $D$ -dim vector.
$\mathbf{X}_{\text{tok}}$	Token embeddings	$[B, S, D]$	$\text{lookup}(\mathbf{E}, \mathbf{T})$ .
$\mathbf{P}$	Positional embedding	$[S, D]$ (or $[B, S, D]$ )	Learned $\mathbf{P}_\theta$ or sinusoidal (fixed); broadcast to $[B, S, D]$ .
$\mathbf{X}_{\text{emb}}$	Sum of token+pos	$[B, S, D]$	$\mathbf{X}_{\text{tok}} + \text{BC}_{B,S}(\mathbf{P})$ .
$\mathbf{X}$	Input to MHA	$[B, S, D]$	After dropout (DO); goes to LN/MHA stack.
$\mathbf{P}_\theta$	Learned pos. params	matches $\mathbf{P}$	Used when positions are trainable; otherwise “sin” denotes fixed sinusoidal.

**Shape symbols:**  $B$ =batch size,  $S$ =sequence length,  $D$ =model dim,  $V$ =vocab size.

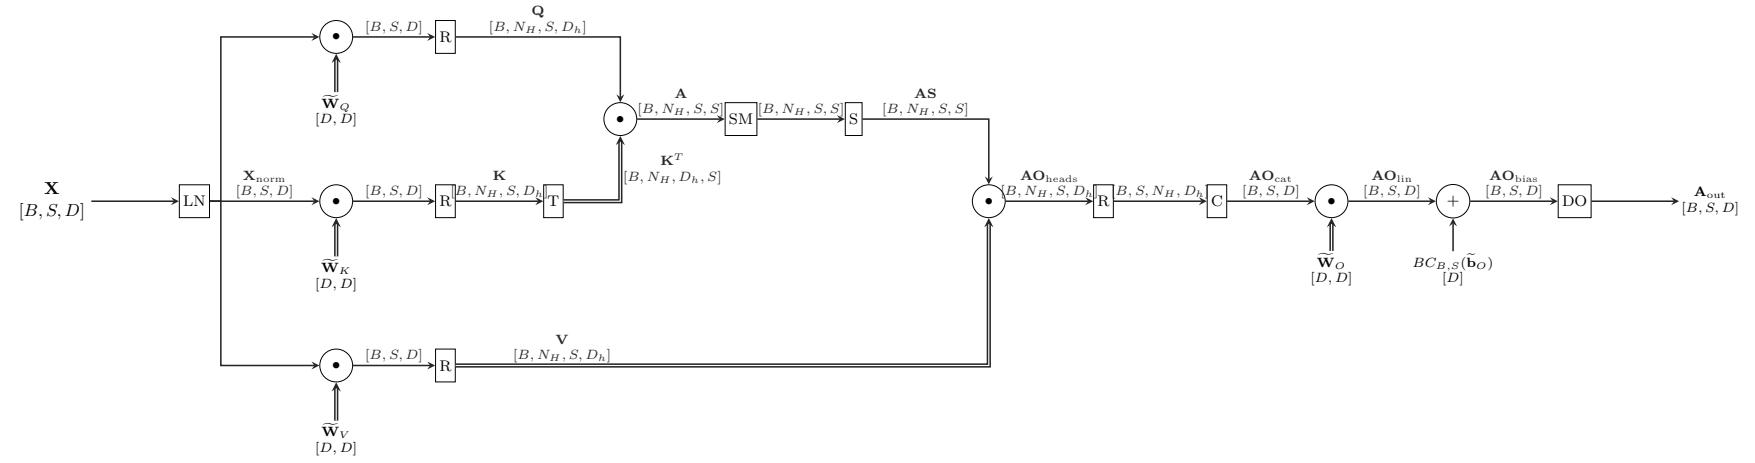
**Notes:** In practice,  $\mathbf{P}$  may be pre-broadcast to  $[B, S, D]$  or added per-token with implicit broadcasting.

### 1.3 Multi-Head Attention (MHA)

Multi-Head Attention enables the model to jointly attend to information from different representation subspaces.

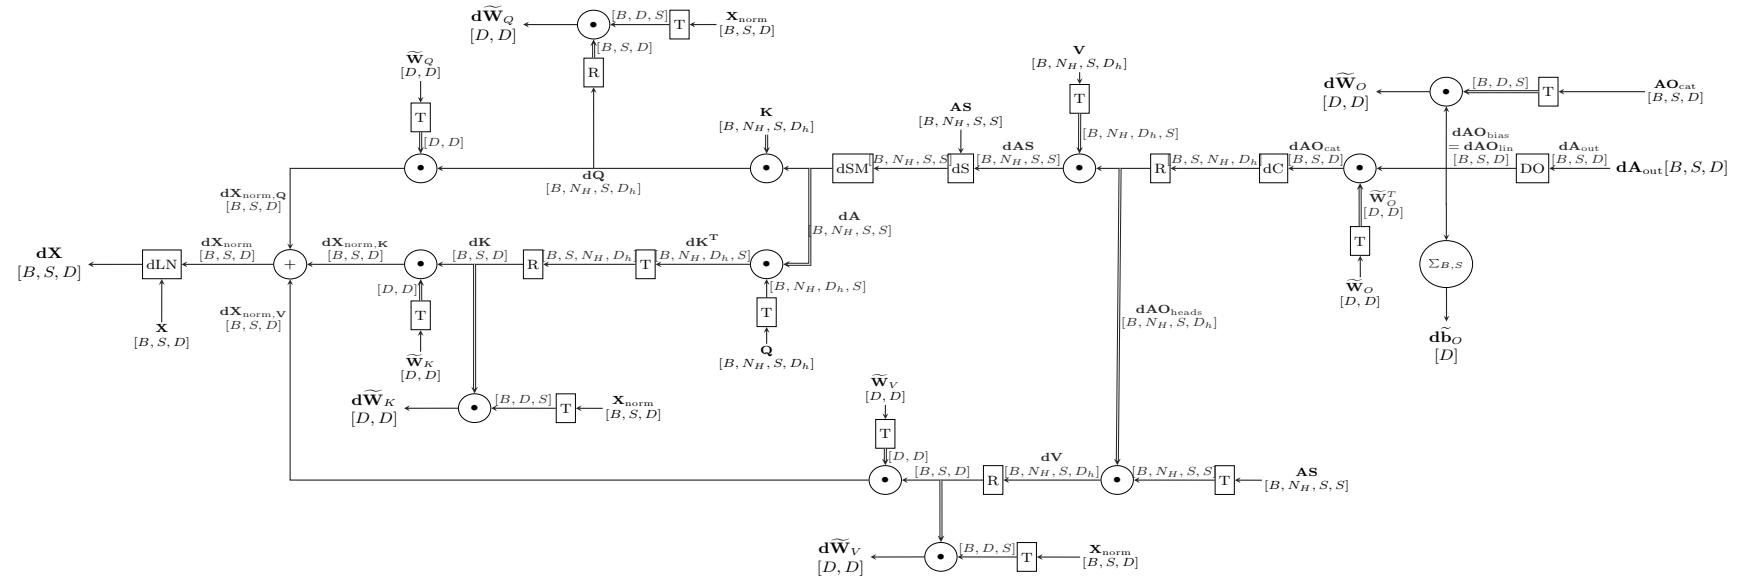
#### 1.3.1 Forward Pass

##### Multi-Head Attention Forward Pass



### 1.3.2 Backward Pass

## Multi-Head Attention Backward Pass





### 1.3.3 Operations Summary

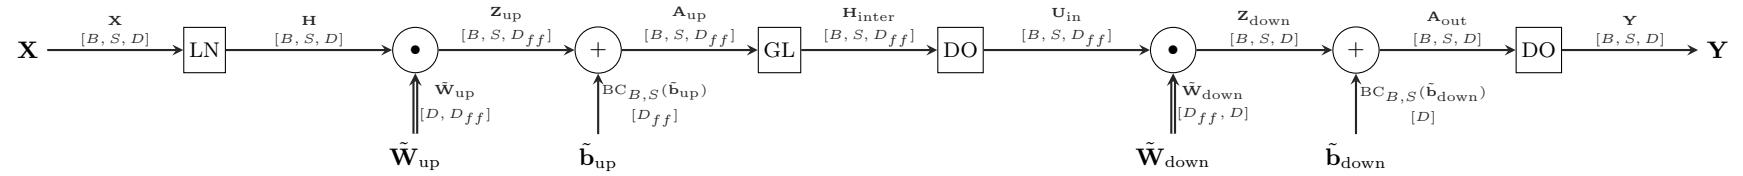
Category	Symbol / Abbrev	Name	Shape / Type	Notes
Ops	LN	Layer Normalization	op	Normalizes per token (last dim $D$ ).
Ops	DO	Dropout	op	Training-time only; identity at inference.
Ops	+	Bias Add	op	Adds broadcast bias; see $\text{BC}_{B,S}(\cdot)$ .
Ops	T	Transpose	op	e.g., $[B, N_H, S, D_h] \rightarrow [B, N_H, D_h, S]$ .
Ops	R	Reshape / Split / Merge	op	$[B, S, D] \leftrightarrow [B, N_H, S, D_h]$ .
Ops	C	Concatenate	op	$[B, S, N_H, D_h] \rightarrow [B, S, D]$ .
Ops	SM	Scale (+ Mask)	op	Multiply by $1/\sqrt{D_h}$ and apply mask.
Ops	S	Softmax	op	Over key length $S$ per head.
Ops	$\text{BC}_{B,S}(\cdot)$	Broadcast	op	Broadcast length- $D$ (or $D_h$ ) to $[B, S, \cdot]$ .
Ops	dS	Softmax Backward	op	Backprop through softmax over $S$ .
Ops	dSM	Scale/Mask Backward	op	Backprop through scaling/masking.
Ops	dC	De-concat (Backward)	op	Split grads from concatenated heads.
Ops	dLN	LayerNorm Backward	op	Uses cached stats $(\mu, \sigma)$ and $\mathbf{X}$ .
Data	$\mathbf{X}$	Input hidden states	$[B, S, D]$	Into MHA block (pre-LN).
Data	$\mathbf{X}_{\text{norm}}$	LN output	$[B, S, D]$	Result of $\text{LN}(\mathbf{X})$ .
Data	$\mathbf{Q}, \mathbf{K}, \mathbf{V}$	Query/Key/Value	$[B, N_H, S, D_h]$	From linear projections of $\mathbf{X}_{\text{norm}}$ .
Data	$\widetilde{\mathbf{W}}_Q$	Q weight	$[D, D]$	Per-head realized via reshape (drawn fused).
Data	$\widetilde{\mathbf{W}}_K$	K weight	$[D, D]$	Same convention.
Data	$\widetilde{\mathbf{W}}_V$	V weight	$[D, D]$	Same convention.
Data	$\widetilde{\mathbf{W}}_O$	Output-proj weight	$[D, D]$	Maps concatenated heads to model dim.
Data	$\tilde{\mathbf{b}}_O$	Output bias	$[D]$	Broadcast via $\text{BC}_{B,S}$ .
Data	$\mathbf{A}$	Attention scores	$[B, N_H, S, S]$	$\mathbf{Q}\mathbf{K}^T/\sqrt{D_h}$ (plus mask).
Data	$\mathbf{AS}$	Attention weights	$[B, N_H, S, S]$	$\text{softmax}(\mathbf{A})$ .
Data	$\mathbf{AO}_{\text{heads}}$	Per-head outputs	$[B, N_H, S, D_h]$	$\mathbf{AS} \cdot \mathbf{V}$ .
Data	$\mathbf{AO}_{\text{cat}}$	Concatenated heads	$[B, S, D]$	After $C$ .
Data	$\mathbf{AO}_{\text{lin}}$	Linear output	$[B, S, D]$	$\mathbf{AO}_{\text{cat}} \widetilde{\mathbf{W}}_O$ .
Data	$\mathbf{AO}_{\text{bias}}$	Bias-added output	$[B, S, D]$	$+ \tilde{\mathbf{b}}_O$ .
Data	$\mathbf{A}_{\text{out}}$	MHA output	$[B, S, D]$	After dropout; to next sublayer.
Data	$d\mathbf{A}_{\text{out}}$	Grad wrt MHA output	$[B, S, D]$	Backprop signal entering MHA.
Data	$d\mathbf{Q}, d\mathbf{K}, d\mathbf{V}$	Gradients for Q/K/V	$[B, N_H, S, D_h]$	From attention-core backward.
Data	$d\mathbf{K}^T$	Grad of $K^T$	$[B, N_H, D_h, S]$	Before transpose/reshape to $d\mathbf{K}$ .
Data	$d\mathbf{AO}_{\text{heads}}$	Grad at heads	$[B, N_H, S, D_h]$	Split from $d\mathbf{AO}_{\text{cat}}$ .
Data	$d\mathbf{X}_{\text{norm}, Q}$	Grad wrt $X_{\text{norm}}$ (Q)	$[B, S, D]$	Via $W_Q^T$ .
Data	$d\mathbf{X}_{\text{norm}, K}$	Grad wrt $X_{\text{norm}}$ (K)	$[B, S, D]$	Via $W_K^T$ .
Data	$d\mathbf{X}_{\text{norm}, V}$	Grad wrt $X_{\text{norm}}$ (V)	$[B, S, D]$	Via $W_V^T$ .
Data	$d\mathbf{X}_{\text{norm}}$	Sum of above	$[B, S, D]$	Input to dLN.
Data	$d\mathbf{X}$	Grad wrt input $X$	$[B, S, D]$	Output of dLN.
Data	$d\widetilde{\mathbf{W}}_Q$	Q weight grad	$[D, D]$	Standard matmul rule.
Data	$d\widetilde{\mathbf{W}}_K$	K weight grad	$[D, D]$	Standard matmul rule.

## 1.4 Feed-Forward Network (FFN/MLP)

The FFN consists of two linear transformations with a non-linear activation function in between.

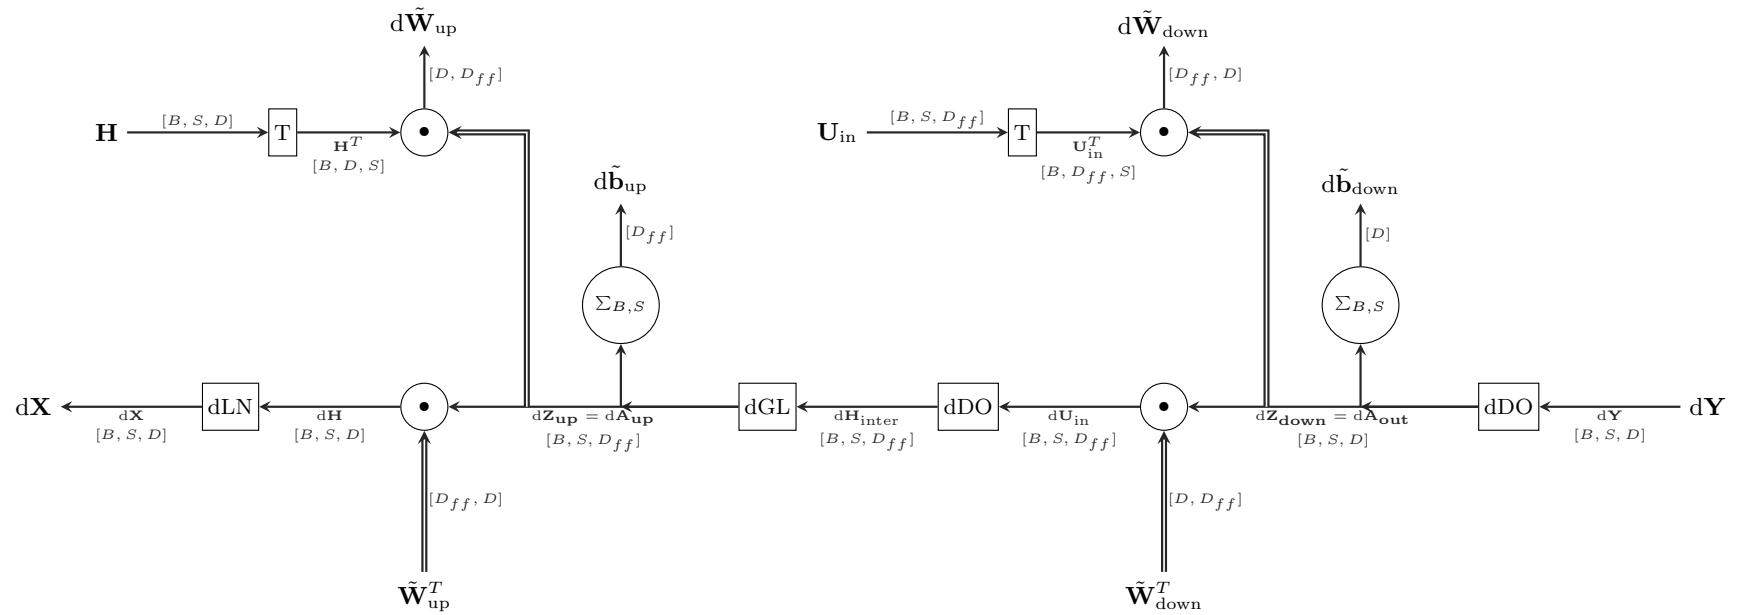
### 1.4.1 Forward Pass

#### MLP Forward Pass



### 1.4.2 Backward Pass

## MLP Backward Pass



### 1.4.3 Operations Summary

MLP (Feed-Forward) Block: Unified Table (Ops & Data)

Category	Symbol / Abbrev	Name	Shape / Type	Notes
Ops	LN	Layer Normalization	op	Normalize per token (last dim $D$ ).
Ops	$\bullet$	Linear (MatMul)	op	Used for up/down projections.
Ops	+	Bias Add	op	Adds broadcast bias; $BC_{B,S}(\cdot)$ .
Ops	GL	GELU (or activation)	op	Nonlinearity on $D_{ff}$ .
Ops	DO	Dropout	op	Training-time only (identity at inference).
Ops	T	Transpose	op	Used in weight-grad computations.
Ops	$\sum_{B,S}$	Reduce-Sum	op	Bias-grad accumulation over batch, seq.
Data	$\mathbf{X}$	Input states	[ $B, S, D$ ]	Block input.
Data	$\mathbf{H}$	LN output	[ $B, S, D$ ]	$LN(\mathbf{X})$ .
Data	$\tilde{\mathbf{W}}_{\text{up}}$	Up weight	[ $D, D_{ff}$ ]	First projection.
Data	$\tilde{\mathbf{b}}_{\text{up}}$	Up bias	[ $D_{ff}$ ]	Broadcast to [ $B, S, D_{ff}$ ].
Data	$\mathbf{Z}_{\text{up}}$	Pre-activation	[ $B, S, D_{ff}$ ]	$H\mathbf{W}_{\text{up}} + \mathbf{b}_{\text{up}}$ .
Data	$\mathbf{A}_{\text{up}}$	Activated	[ $B, S, D_{ff}$ ]	$f(\mathbf{Z}_{\text{up}})$ .
Data	$\mathbf{H}_{\text{inter}}$	Post-DO	[ $B, S, D_{ff}$ ]	After Dropout.
Data	$\tilde{\mathbf{W}}_{\text{down}}$	Down weight	[ $D_{ff}, D$ ]	Second projection.
Data	$\tilde{\mathbf{b}}_{\text{down}}$	Down bias	[ $D$ ]	Broadcast to [ $B, S, D$ ].
Data	$\mathbf{Z}_{\text{down}}$	Linear output	[ $B, S, D$ ]	$H_{\text{inter}}\mathbf{W}_{\text{down}} + \mathbf{b}_{\text{down}}$ .
Data	$\mathbf{A}_{\text{out}}$	Bias-added	[ $B, S, D$ ]	Before dropout (out).
Data	$\mathbf{Y}$	Block output	[ $B, S, D$ ]	After Dropout.
Data	$d\mathbf{Y}$	Grad output	[ $B, S, D$ ]	Incoming grad.
Data	$d\mathbf{Z}_{\text{down}}$	Grad lin-out	[ $B, S, D$ ]	Equals $d\mathbf{A}_{\text{out}}$ .
Data	$d\mathbf{U}_{\text{in}}$	Grad into down	[ $B, S, D_{ff}$ ]	To weight/bias grads.
Data	$d\mathbf{Z}_{\text{up}}$	Grad pre-act	[ $B, S, D_{ff}$ ]	Equals $d\mathbf{A}_{\text{up}} \cdot f'$ .
Data	$d\mathbf{H}$	Grad LN out	[ $B, S, D$ ]	Into dLN.
Data	$d\mathbf{X}$	Grad input	[ $B, S, D$ ]	Block input grad.
Data	$d\tilde{\mathbf{W}}_{\text{up}}$	Weight grad up	[ $D, D_{ff}$ ]	From $H^T$ and $d\mathbf{Z}_{\text{up}}$ .
Data	$d\tilde{\mathbf{W}}_{\text{down}}$	Weight grad down	[ $D_{ff}, D$ ]	From $U_{\text{in}}^T$ and $d\mathbf{Z}_{\text{down}}$ .
Data	$d\tilde{\mathbf{b}}_{\text{up}}$	Bias grad up	[ $D_{ff}$ ]	$\sum_{B,S} d\mathbf{Z}_{\text{up}}$ .
Data	$d\tilde{\mathbf{b}}_{\text{down}}$	Bias grad down	[ $D$ ]	$\sum_{B,S} d\mathbf{Z}_{\text{down}}$ .

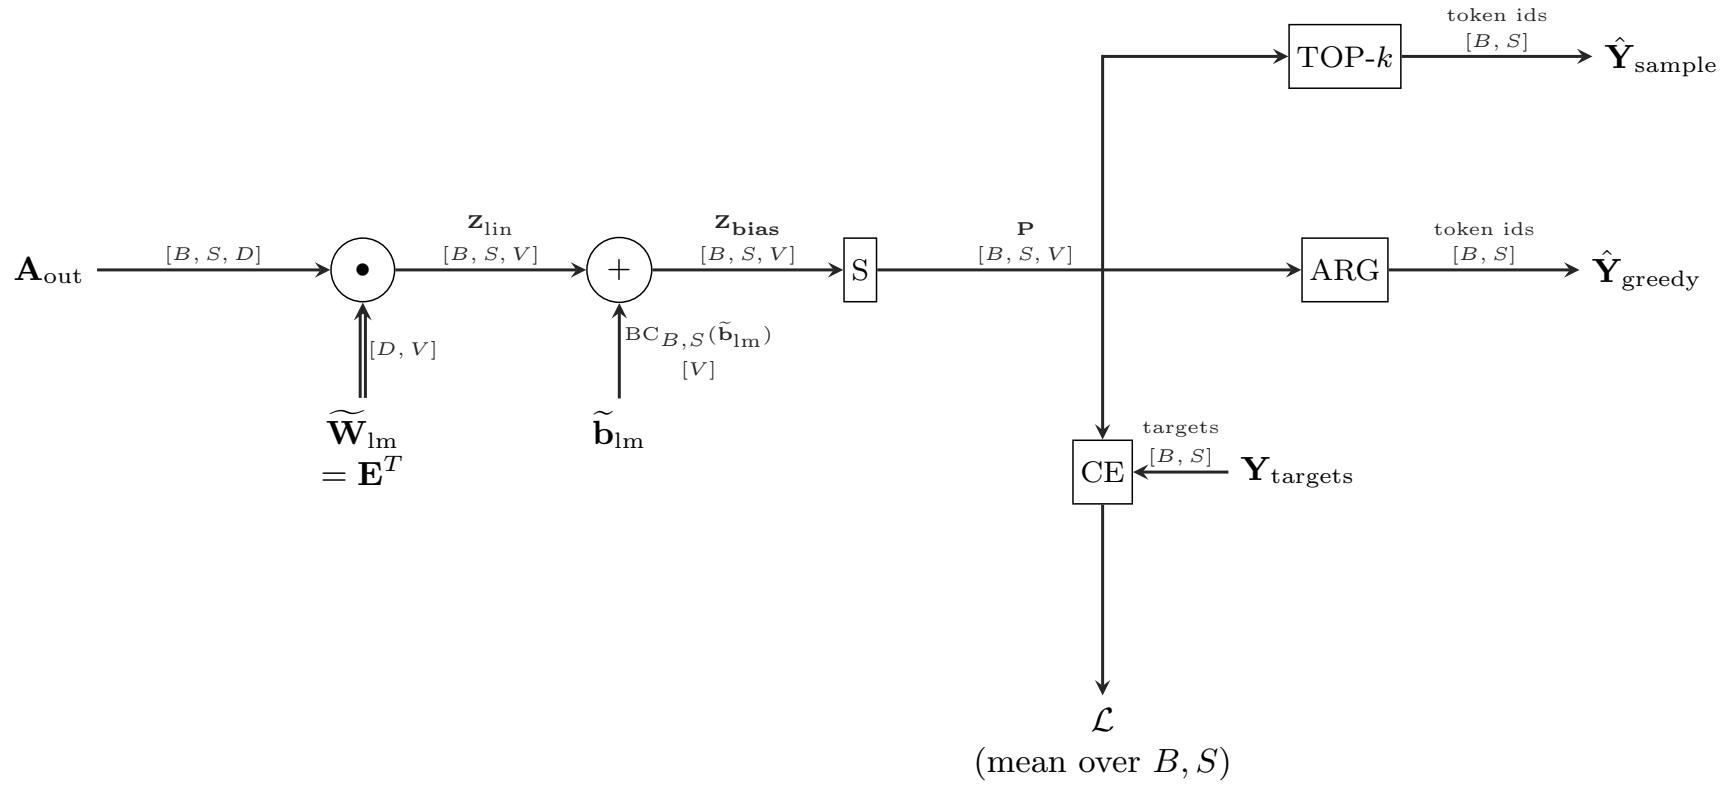
**Shape symbols:**  $B$ =batch,  $S$ =sequence,  $D$ =model dim,  $D_{ff}$ =FFN hidden dim (e.g.,  $4 \times D$ ).

## 1.5 Output Projection & Loss Computation

The final layer projects the hidden states to vocabulary logits and computes the cross-entropy loss.

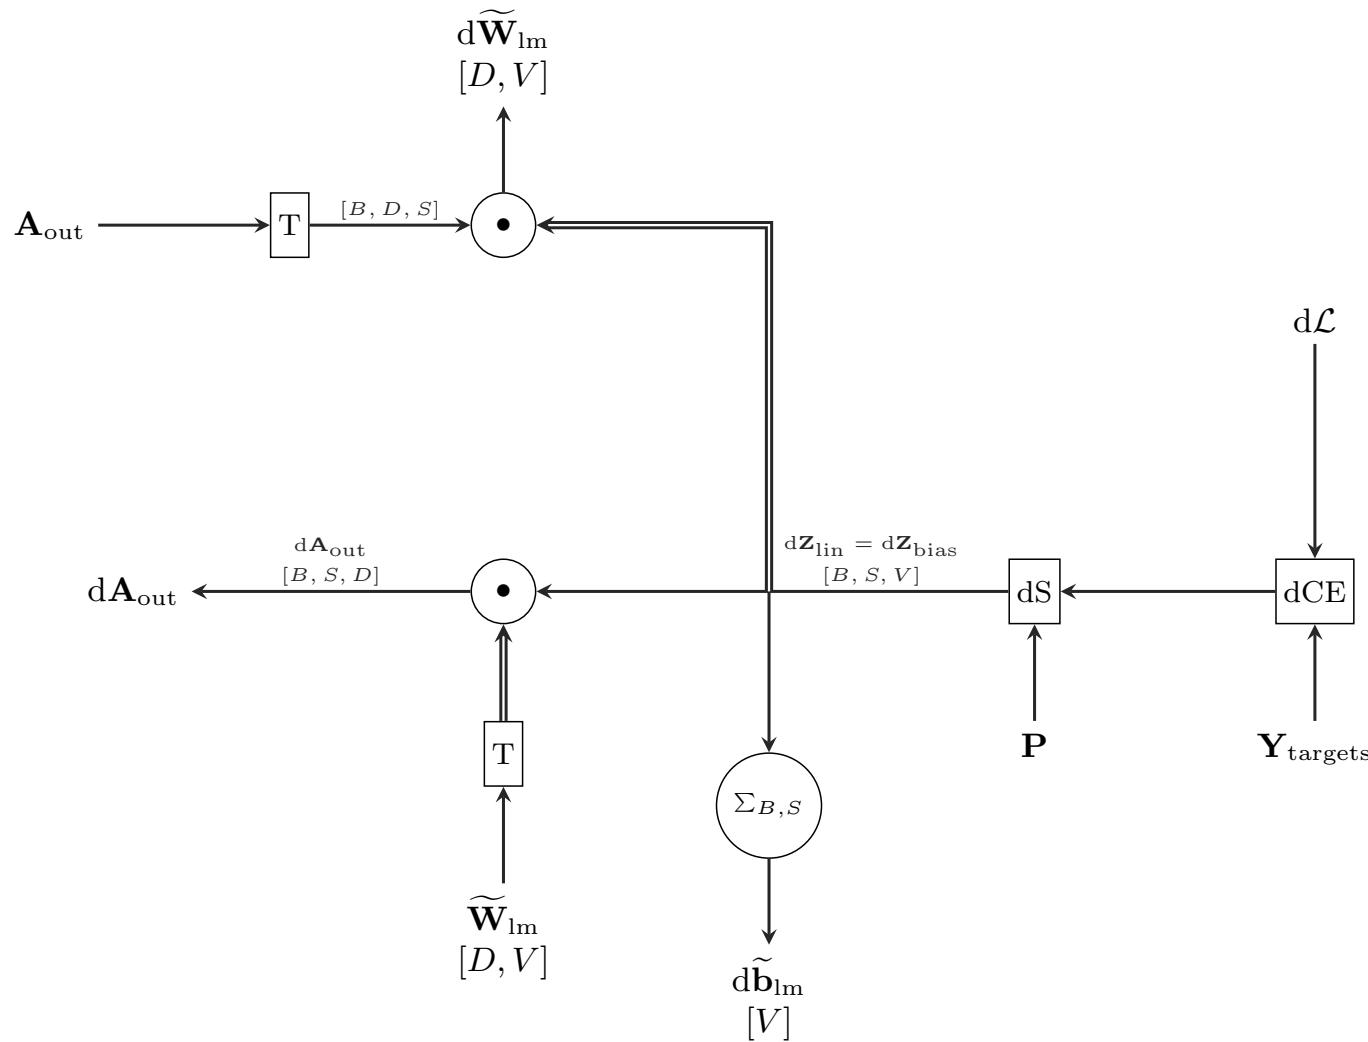
### 1.5.1 Forward Pass

### Token Generation & Loss (Forward)



### 1.5.2 Backward Pass

## Token Generation & Loss — Backward (Corrected)



### 1.5.3 Operations Summary

	Abbrev	Name	Type / Shape	Notes
Operations (Ops)	S	Softmax	op	Over vocab axis $V$ ; outputs probabilities $\mathbf{P}$ .
	CE	Cross-Entropy	op	Usually <i>sparse</i> CE consuming label indices $\mathbf{Y}$ .
	ARG	Argmax (greedy)	op	$\text{argmax}_V$ to get token ids (no gradient).
	TOP- $k$	Top- $k$ / sampling	op	Optional decoding path; no gradient.
	T	Transpose	op	E.g., $\widetilde{\mathbf{W}}_{\text{lm}}^T \in \mathbb{R}^{V \times D}$ .
	$\text{BC}_{B,S}(\cdot)$	Broadcast	op	Expand $[V] \rightarrow [B, S, V]$ for bias add.
	dS	Softmax backward	op	With CE: $d\mathbf{Z}_{\text{bias}} = \mathbf{P} - \text{onehot}(\mathbf{Y})$ .
	dAddB	Addition (Bias) backward	op	Sends $d\mathbf{Z}_{\text{bias}}$ to matmul and $\sum_{B,S}$ .
	$\sum_{B,S}$	Summation	op	Yields $d\tilde{\mathbf{b}}_{\text{lm}}$ .

Data Tensors (Values)

Symbol	Name	Shape	Notes
$\mathbf{A}_{\text{out}}$	Transformer output (hidden)	$[B, S, D]$	Final hidden from the Transformer block(s).
$\widetilde{\mathbf{W}}_{\text{lm}}$	LM head weight (tied)	$[D, V]$	Typically tied to $\mathbf{E}^T$ .
$\tilde{\mathbf{b}}_{\text{lm}}$	LM head bias	$[V]$	Broadcast-added over $[B, S, V]$ .
$\mathbf{Z}_{\text{lin}}$	Logits (linear output)	$[B, S, V]$	$\mathbf{A}_{\text{out}} \widetilde{\mathbf{W}}_{\text{lm}}$ .
$\mathbf{Z}_{\text{bias}}$	Logits (final/Softmax input)	$[B, S, V]$	$\mathbf{Z}_{\text{lin}} + \tilde{\mathbf{b}}_{\text{lm}}$ .
$\mathbf{P}$	Probabilities	$[B, S, V]$	$\text{softmax}(\mathbf{Z}_{\text{bias}})$ .
$\mathbf{Y}$	Target token ids	$[B, S]$	Ground-truth indices (sparse labels).
$\mathcal{L}$	Loss	scalar or $[B, S]$	Typically mean over $B, S$ .
$d\mathcal{L}$	Loss gradient	scalar-grad	Starting signal for backward pass.
$d\mathbf{Z}_{\text{bias}}$	Final Logits gradient	$[B, S, V]$	From CE+Softmax: $\mathbf{P} - \text{onehot}(\mathbf{Y})$ .
$d\mathbf{Z}_{\text{lin}}$	Linear output grad	$[B, S, V]$	Same as $d\mathbf{Z}_{\text{bias}}$ .
$d\widetilde{\mathbf{W}}_{\text{lm}}$	LM weight grad	$[D, V]$	$= \mathbf{A}_{\text{out}}^T d\mathbf{Z}_{\text{lin}}$ .
$d\tilde{\mathbf{b}}_{\text{lm}}$	LM bias grad	$[V]$	$= \sum_{B,S} d\mathbf{Z}_{\text{bias}}$ .
$d\mathbf{A}_{\text{out}}$	Hidden grad	$[B, S, D]$	$= d\mathbf{Z}_{\text{lin}} \widetilde{\mathbf{W}}_{\text{lm}}^T$ .

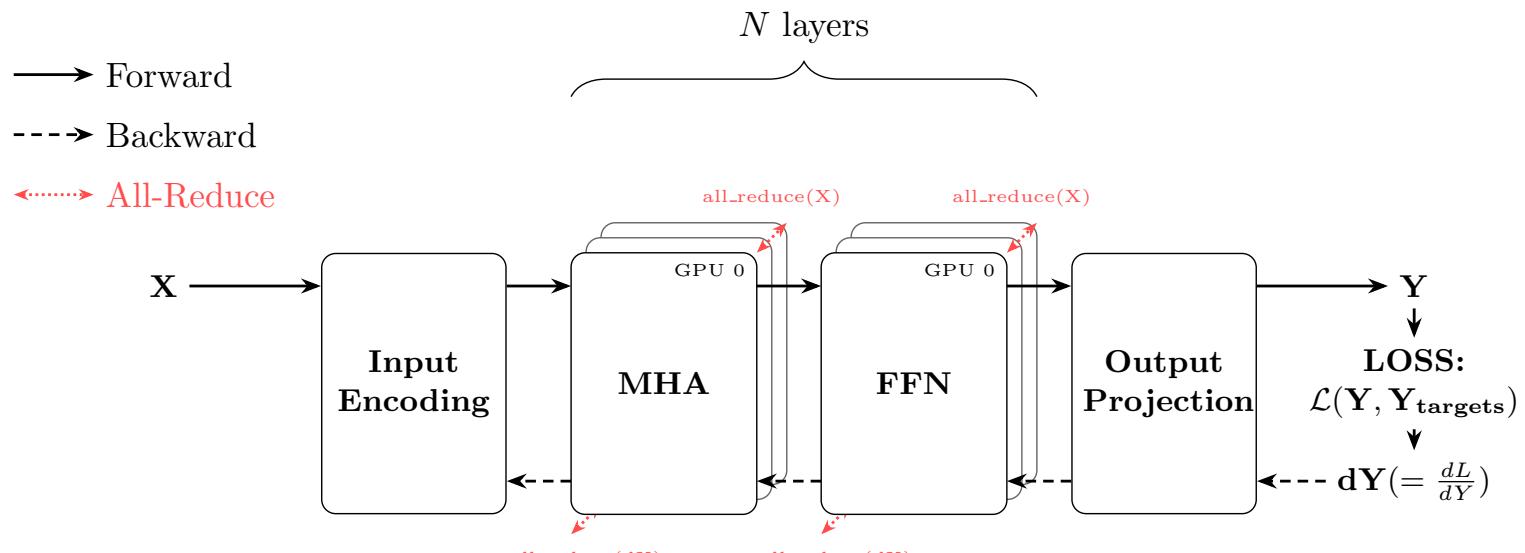
Shapes:  $B$ =batch,  $S$ =sequence length,  $D$ =hidden dim,  $V$ =vocab size.

## 2 Tensor Parallelism (TP)

*[This chapter will cover Tensor Parallelism implementation details]*

## 2.1 TP Overview

# Transformer Overall Flow (TP with 3 GPUs)



[To be added]

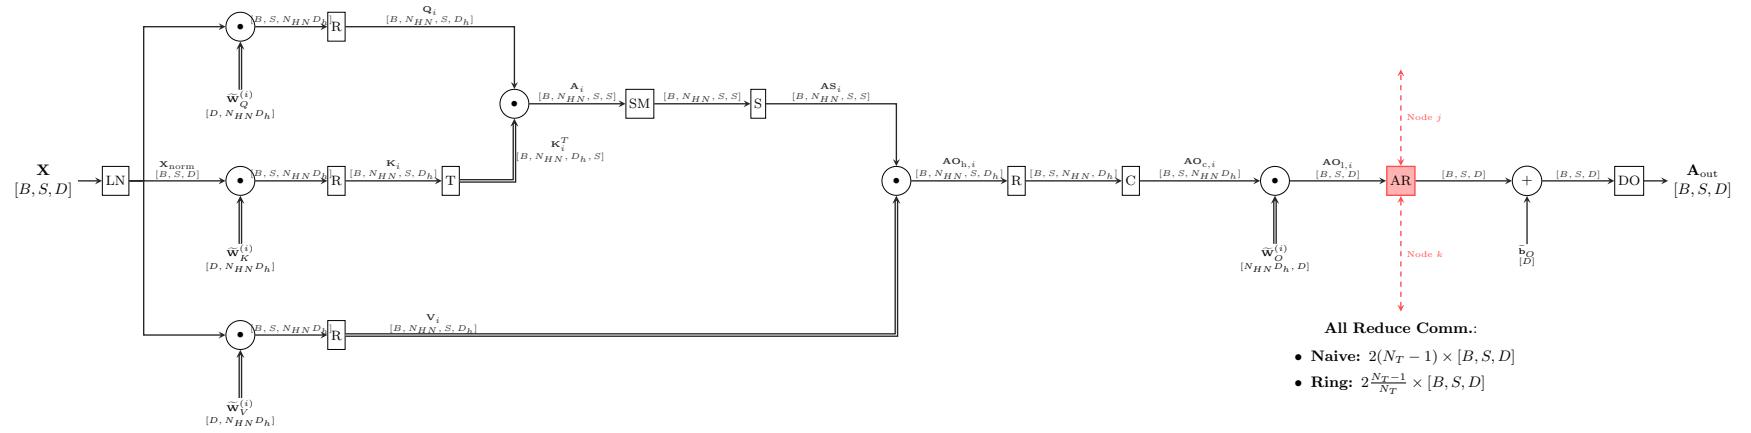
## 2.2 Forward Pass

In the tensor-parallel setting, the multi-head attention is distributed across  $N_T$  devices (GPUs). The notation used in the diagram is defined as follows:

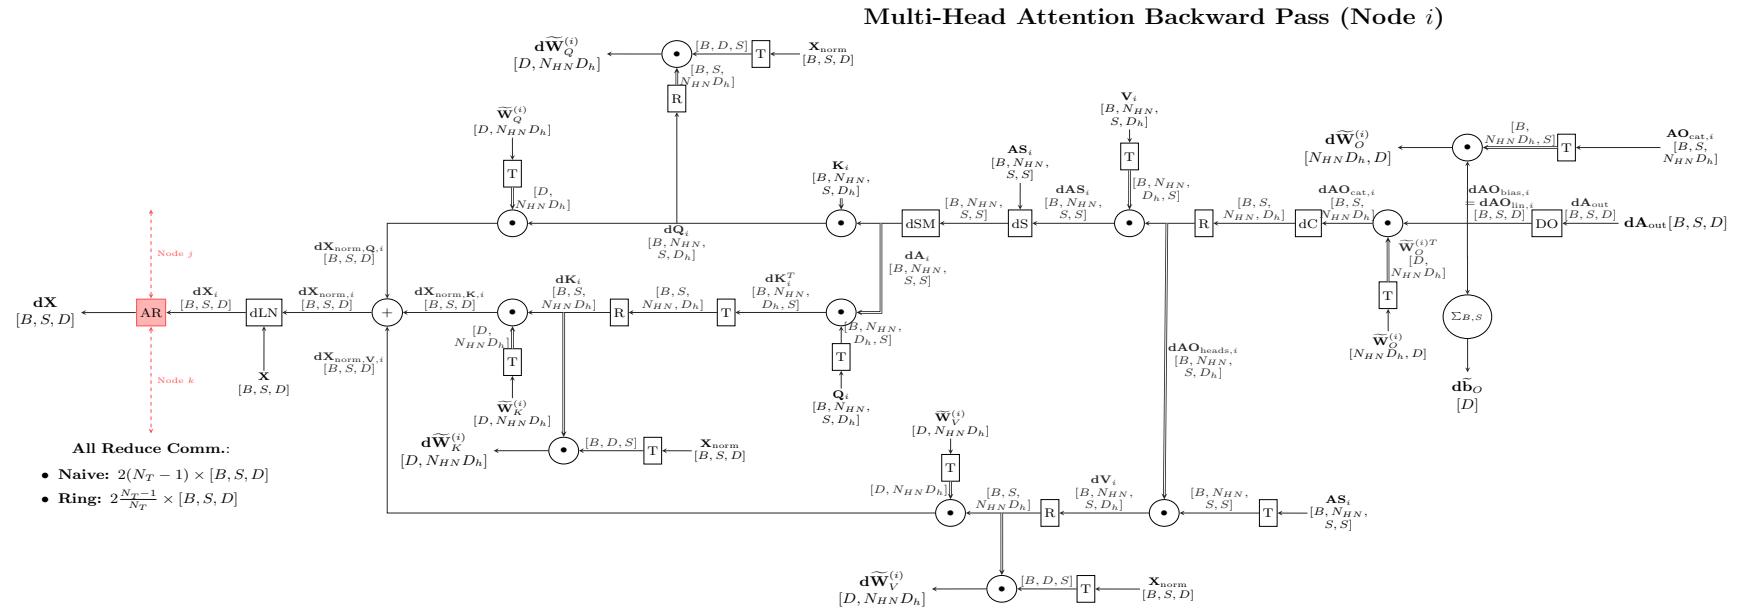
- $N_H$ : Total number of attention heads (global)
- $N_T$ : Tensor parallelism degree (number of devices)
- $N_{HN}$ : Number of heads per device (local), where  $N_{HN} = \frac{N_H}{N_T}$

Each device processes  $N_{HN}$  heads independently. The weight matrices  $\tilde{\mathbf{W}}_Q$ ,  $\tilde{\mathbf{W}}_K$ ,  $\tilde{\mathbf{W}}_V$ , and  $\tilde{\mathbf{W}}_O$  are column-partitioned across devices, with each device holding a  $[D, N_{HN} \cdot D_h]$  slice of the original  $[D, D]$  matrix.

Multi-Head Attention Forward Pass (Node  $i$ )



## 2.3 Backward Pass



## 2.4 Communication Patterns

[To be added]



### 3 Data Parallelism (DP)

#### 3.1 DP Overview

